

# Selected Topics in Network Optimization: Aligning BDDs for a Facility Location Problem and a Search Tree Method for DSPI.

Dissertation defense


Alexey Bochkarev

2021-11-02, 08-30, Zoom / Freeman Hall 123

# Outline

- 1 **TODO Aligning BDDs**
  - On BDD representations
  - BDD alignment (“original”) problem
  - The simplified problem
  - Numerical experiments
- 2 **TODO A BDD-based approach to a Facility Location Problem**
  - Problem description
  - BDD representation
  - Solving the CPP
  - Numerical experiments
- 3 **Monte Carlo Tree Search for DSPI**
  - Problem formulation
  - What do we propose?
  - MCTS framework
  - Numerical experiments
  - On correctness
- 4 **TODO Conclusion**
  - Summary
  - Future research

# A game of "interdiction": intro



- **Network:** a directed graph with two special nodes (source  $\textcircled{s}$  and terminal  $\textcircled{t}$ ), and a pair of "costs" associated to each edge.
- **User:** seeks to run through the graph,  $\textcircled{s}$  to  $\textcircled{t}$ , at min cost.
- **Attacker:** maximizes the User's cost by "attacking" the arcs, having a limited "budget".

We consider a **dynamic** version of the game, following [Sefair and Smith(2016)].  
(NP-hard)

# A game of “interdiction”: formulation

The Interdictor's optimal objective  $z^*$  can be expressed as:

$$z^*(S, i) = \max_{S' \subseteq \text{FS}(i) \setminus S : |S \cup S'| \leq b} \left\{ \min_{j \in \text{FS}(i)} \{z^*(S \cup S', j) + \tilde{c}_{ij}(S \cup S')\} \right\},$$

where:

- $S$ : interdiction set,
- $i$ : current Evader's node,  $\text{FS}(i)$  – forward star of node  $i$ ,
- $\tilde{c}_{ij}$ : arc traversal costs (given the interdiction),
- $b$ : Interdictor's budget.

# A game of “interdiction”: formulation

The Interdictor's optimal objective  $z^*$  can be expressed as:

$$z^*(S, i) = \max_{S' \subseteq \text{FS}(i) \setminus S : |S \cup S'| \leq b} \left\{ \min_{j \in \text{FS}(i)} \{z^*(S \cup S', j) + \tilde{c}_{ij}(S \cup S')\} \right\},$$

where:

- $S$ : interdiction set,
- $i$ : current Evader's node,  $\text{FS}(i)$  – forward star of node  $i$ ,
- $\tilde{c}_{ij}$ : arc traversal costs (given the interdiction),
- $b$ : Interdictor's budget.

## Existing algorithms (by Sefair & Smith)

- polynomial DP algorithm for a DAG
- exact DP algorithm, exp time for general case.

What do we propose?

# Monte Carlo Tree Search

- Maintain the game tree,
- Try not to create all the nodes,
- Prune the definitely suboptimal ones,
- Drive the tree growth by a computationally cheap objective estimate (e.g., based on simulated games).

# The “game tree”

Create a “game tree”, where **nodes** contain the following information.

- Current **status**
  - $p(j) \in \mathcal{N}$ : where is the Evader,
  - $S_j \subseteq \mathcal{A}$ : what is interdicted,
  - $\tau(j)$ : who's turn is it (Interdictor/Evader)

# The “game tree”

Create a “game tree”, where **nodes** contain the following information.

- Current **status**
  - $p(j) \in \mathcal{N}$ : where is the Evader,
  - $S_j \subseteq \mathcal{A}$ : what is interdicted,
  - $\tau(j)$ : who's turn is it (Interdictor/Evader)
- Possible further **development**
  - $\text{children}(j)$ : child game tree nodes,
  - $\text{actions}(j)$ : available actions.



# The “game tree”

Create a “game tree”, where **nodes** contain the following information.

- Current **status**
  - $p(j) \in \mathcal{N}$ : where is the Evader,
  - $S_j \subseteq \mathcal{A}$ : what is interdicted,
  - $\tau(j)$ : who's turn is it (Interdictor/Evader)
- Possible further **development**
  - $\text{children}(j)$ : child game tree nodes,
  - $\text{actions}(j)$ : available actions.
- **Costs** info, to drive the search and prune the tree
  - $\hat{Q}_j$ : cost-to-go (starting from this node),
  - $LB(j)$ ,  $UB(j)$ : bounds on the true cost-to-go,
  - $N_j \in \mathbb{N}$ : how many times the node was visited


# The “game tree”

Create a “game tree”, where **nodes** contain the following information.

- Current **status**
  - $p(j) \in \mathcal{N}$ : where is the Evader,
  - $S_j \subseteq \mathcal{A}$ : what is interdicted,
  - $\tau(j)$ : who's turn is it (Interdictor/Evader)
- Possible further **development**
  - $\text{children}(j)$ : child game tree nodes,
  - $\text{actions}(j)$ : available actions.
- **Costs** info, to drive the search and prune the tree
  - $\hat{Q}_j$ : cost-to-go (starting from this node),
  - $LB(j)$ ,  $UB(j)$ : bounds on the true cost-to-go,
  - $N_j \in \mathbb{N}$ : how many times the node was visited

So, we iterate through **episodes**, each one implying a “full cycle” of the game tree update in four **phases**.

# Phase 1. Selection




## What's happening:

- Start at the root node,
- Use *tree policy* to choose the next node recursively...
- ... pruning nodes as we go, when possible ...
- ... until we reach a leaf.

## What's updated:

- Nothing in the tree.
- Along the way: bounds for pruning (more momentarily!) + path costs.

# Phase 2. Expansion



## What's happening:


- Create child nodes for possible actions.

## What's updated:

- New nodes are created,
- UBs and LBs are calculated

**Note:** Some inconsistencies can be introduced here, between child and parent nodes.

# Phase 3. Roll-outs



## What's happening:


- Run a random simulated game from each node,
- Calculate cost-to-go estimate  $\hat{Q}_j$  as the simulated game cost.

## What's updated

- Cost-to-go for each new node.

**Note:** We do not record the intermediate game states occurred during roll-outs!

# Phase 4. Backpropagation



## What's happening:

- Start at the selected node,
- Recursively update (“propagate”) node information for parents ...
- ... until we reach the root.

**What's updated:** Information in each parent node, using the child nodes:

- UBs and LBs
- Cost-to-go estimate: the best value (given the turn).

# The Algorithm

The algorithm can perform actions for both players. Each turn involves two steps:

## Step 1. THINK.

We iteratively improve the tree (while we have budget):

**FOR**  $k = 1, \dots, K$  **DO**

- Selection
- Expansion
- Roll-outs
- Backpropagation

**END.**

## Step 2. ACT.

... then pick an action corresponding to the “most attractive” child node of the root.

# There are several secret ingredients





# SI-1. How to select?

- **with probability**  $\varepsilon$ : choose at random;
- **otherwise**, a child node with the **best score**:

$$R_j = \underbrace{\sigma_i(\tilde{C}_{ij} + \hat{Q}_j)}_{\text{best cost-to-go}} + \underbrace{C_p \sqrt{\log(N_i)/N_j}}_{\text{encourage exploration}}, \quad \text{for all } j \in \text{children}(i)$$

“Best” here depends on the turn (the Evader will choose the smallest cost estimate, the Interdictor — the largest).

# SI-2. How to prune?


We leverage the classic idea of **alpha-beta pruning**:

Maintain two running numbers (bounds):

- $\alpha$ : the worst (minimum) alternative cost achievable by the Interdictor,
- $\beta$ : the worst (maximum) alternative cost achievable by the Evader.

**Pruning condition:**  $\beta \leq \hat{\alpha}_j$  or  $\hat{\beta}_j \leq \alpha$ , where

- $\hat{\alpha}_j = \pi_n + LB(j)$  (Interdictor's turns),  
and
- $\hat{\beta}_j = \pi_n + \tilde{C}_{nj} + UB(j)$  (Evader's turns)



## SI-3. How to back-propagate?

Assuming the Evader's turn, and  $i$  being the current game tree node:

- Update the bounds:

$$UB(i) \leftarrow \min_{j \in \text{children}(i)} \left\{ \tilde{C}_{ij}(S_i) + UB(j) \right\},$$

$$LB(i) \leftarrow \min_{j \in \text{children}(i)} \left\{ \tilde{C}_{ij}(S_i) + LB(j) \right\}.$$


- Update the cost-to-go estimate:

$$\hat{Q}_i \leftarrow \min_{j \in \text{children}(i)} \left\{ \tilde{C}_{ij}(S_i) + \hat{Q}_j \right\}.$$

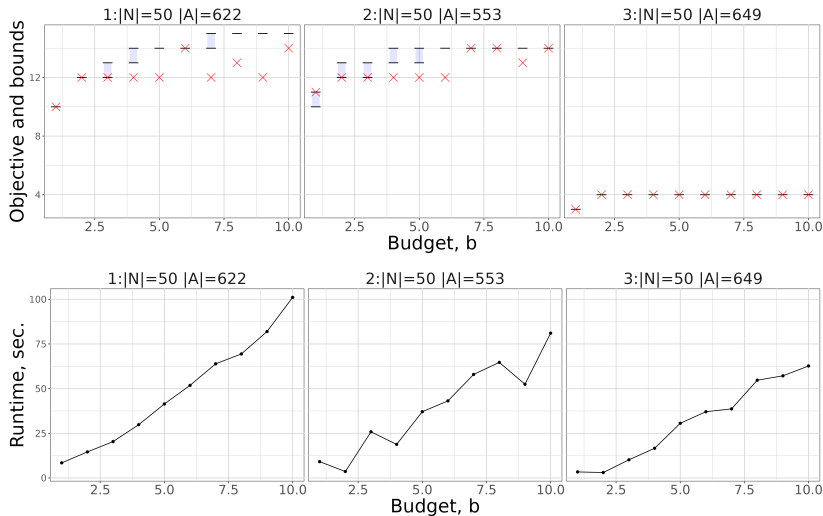
# Numerical experiments: strategy

- How does it perform on pre-defined instances? (relative to the known optimum, and to the bounds)
- How does it perform on randomly generated instances with different budgets?
- What's the dynamics of the tree construction? How does the algorithm work?
- What's the point of “playing out”, i.e., changing root nodes?

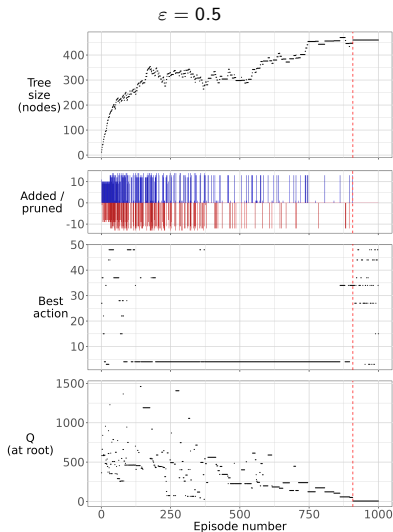
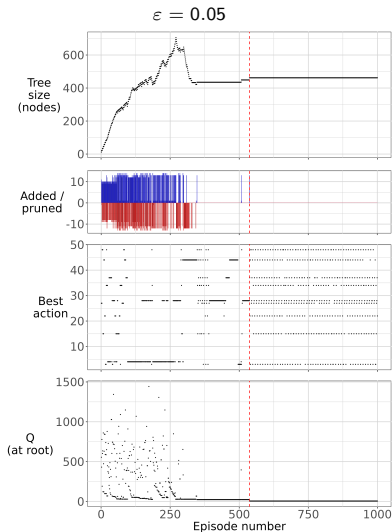
# Pre-defined instances




# Random instances (snapshot)



# Convergence profiles







# A remark: that's not just different runs



For each value of  $\epsilon$  (0.05, 0.5, and 1) we performed three runs, to confirm these are indeed different “modes” of the algorithm.



# Play-out vs. first-move strategy



Objective estimate (relative to the optimum)

Runtime per instance (seconds)

# Why does it work? (A sketch on “correctness”)

## Theorem (Proposition)


$$\lim_{k \rightarrow \infty} \mathbb{P}\{Q_{root}^k = \text{true optimum}\} = 1$$

A sketch of the proof:


- The game tree has finite number of nodes (there is a bound independent from  $K$ ).
- We never cut off all the optima  $\Rightarrow$  the tree contains at least one.
- What is left is a finite-size minimax tree, containing an optimum.
- As  $K \rightarrow \infty$ , probability to select every node for expansion converges to 1.

# Why in the world the tree is finite?

Network:




After the **first** expansion:





# Why in the world the tree is finite?

After the **second** expansion:

Network:



After the **first** expansion:



# Mentioned sources



Jorge A. Sefair and J. Cole Smith.

Dynamic shortest-path interdiction.

*Networks*, 68(4):315–330, 2016.

ISSN 1097-0037.

doi: 10.1002/net.21712.